References

Antoine Bosselut, Over Levy, Ari Holtzman, Coin Ennist, Dieter Fox, and Yejin Choi. 2018. Simulating action dynamics with neural process networks. In *ICLR*.

Guan-Lin Chao and Ian Lane. 2019. Bert-dst: Scalable end-to-end dialogue state tracking with bidirectional encoder representations from transformer. In *Interspeech*.

Angela Fan, Mike Lewis, and Yann Dauphin. 2018. Hierarchical neural story generation. In *ACL*.

Angela Fan, Mike Lewis, and Yann Dauphin. 2019. Strategies for structuring story generation. In *ACL*.

Marjan Ghazvininejad, Xing Shi, Yejin Choi, and Kevin Knight. 2016. Generating topical poetry. In *EMNLP*.

Brent Harrison, Christopher Purdy, and Mark O. Riedl. 2017. Toward automated story generation with markov chain monte carlo methods and deep neural networks.

Mikael Henaff, Jason Weston, Arthur Szlam, Antoine Bordes, and Yann LeCun. 2017. Tracking the world state with recurrent entity networks. In *ICLR*.

Ari Holtzman, Jan Buys, Maxwell Forbes, and Yejin Choi. 2019. The curious case of neural text degeneration. *ArXiv*, abs/1904.09751.

Daphne Ippolito, David Grangier, Chris Callison-Burch, and Douglas Eck. 2019. Unsupervised hierarchical story infilling. In *Proceedings of the First Workshop on Narrative Understanding*, pages 37–43, Minneapolis, Minnesota. Association for Computational Linguistics.

Parag Jain, Priyanka Agrawal, Abhijit Mishra, Mohak Sukhwani, Anirban Laha, and Karthik Sankaranarayanan. 2017. Story generation from sequence of independent short descriptions. *ArXiv*, abs/1707.05501.

Nitish Shirish Keskar, Bryan McCann, Lav R Varshney, Caiming Xiong, and Richard Socher. 2019. Ctrl: A conditional transformer language model for controllable generation. *arXiv preprint arXiv:1909.05858*.

Chloe Kiddon, Luke Zettlemoyer, and Yejin Choi. 2016. Globally coherent text generation with neural checklist models. In *EMNLP*.

Michael Lebowitz. 1987. Planning stories. In *Proceedings of the 9th annual conference of the cognitive science society*, pages 234–242.

Sungjin Lee. 2013. Structured discriminative model for dialog state tracking. In *SIGDIAL*.

Boyang Li, Stephen Lee-Urban, George Johnston, and Mark O. Riedl. 2013. Story generation with crowdsourced plot graphs. In *AAAI*.

Jiwei Li, Thang Luong, and Dan Jurafsky. 2015. A hierarchical neural autoencoder for paragraphs and documents. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1106–1115, Beijing, China. Association for Computational Linguistics.

Chin-Yew Lin. 2004. ROUGE: A package for automatic evaluation of summaries. In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics. https://pypi.org/project/rouge/.

Lara J. Martin, Prithviraj Ammanabrolu, William Hancock, S. Singh, Brent Harrison, and Mark O. Riedl. 2017. Event representations for automated story generation with deep neural nets. In *AAAI*.

Nanyun Peng, Marjan Ghazvininejad, Jonathan May, and Kevin Knight. 2018. Towards controllable story generation. In *Proceedings of the First Workshop on Storytelling*.

Rafael Pérez y Pérez and Mike Sharples. 2001. Mexica: A computer model of a cognitive account of creative writing. *Journal of Experimental & Theoretical Artificial Intelligence*, 13:119 – 139.

Julie Porteous and Marc Cavazza. 2009. Controlling narrative generation with planning trajectories: the role of constraints. In *Joint International Conference on Interactive Digital Storytelling*.

Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre-training.

Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners.

Mark O Riedl. 2009. Story planning: Creativity through exploration, retrieval, and analogical transformation. In *Minds and Machines*, volume 20(4):589–614.

Mark O Riedl and Robert Michael Young. 2010. Narrative planning: Balancing plot and character. In *Journal of Artificial Intelligence Research*.

Stuart Rose, Nick Cramer, and Dave Engel. 2010. Automatic keyword extraction from individual documents. *Text Mining: Applications and Theory edited by Michael W. Berry and Jacob Kogan, John Wiley & Sons, Ltd*.

Evan Sandhaus. 2008. The new york times annotated corpus ldc2008t19. *DVD. Philadelphia: Linguistic Data Consortium*.

Abigail See, Aneesh Pappu, Rohun Saxena, Akhila Yerukola, and Christopher D. Manning. 2019. Do massively pretrained language models make better storytellers? In *CoNLL*, volume abs/1909.10705.

Dinghan Shen, Asli Celikyilmaz, Yizhe Zhang, Liqun Chen, Xin Wang, Jianfeng Gao, and Lawrence Carin. 2019. Towards generating long and coherent text with multi-level latent variable models. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2079–2089, Florence, Italy. Association for Computational Linguistics.

Sainbayar Sukhbaatar, Arthur Szlam, Jason Weston, and Rob Fergus. 2015. End-to-end memory networks. In *NeurIPS*, pages 2440–2448.

Pradyumna Tambwekar, Murtaza Dhuliawala, Lara J. Martin, Animesh Mehta, Brent Harrison, and Mark O. Riedl. 2019. Controllable neural story plot generation via reward shaping. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19*, pages 5982–5988. International Joint Conferences on Artificial Intelligence Organization.

Blaise Thomson and Steve J Young. 2010. Bayesian update of dialogue state: A pomdp framework for spoken dialogue systems. *Computer Speech and Languag*, 24(4):562–588.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *NeurIPS*, pages 5998–6008.

Zhe Wang, Wei He, Hua Wu, Haiyang Wu, Wei Li, Haifeng Wang, and Enhong Chen. 2016. Chinese poetry generation with planning based neural network. In *Coling*.

Tsung-Hsien Wen, Milica Gašić, Nikola Mrkšić, Pei-Hao Su, David Vandyke, and Steve Young. 2015. Semantically conditioned LSTM-based natural language generation for spoken dialogue systems. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 1711–1721, Lisbon, Portugal. Association for Computational Linguistics.

Jingjing Xu, Xuancheng Ren, Yi Zhang, Qi Zeng, Xiaoyan Cai, and Xu Sun. 2018. A skeleton-based model for promoting coherence among sentences in narrative story generation. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 4306–4315, Brussels, Belgium. Association for Computational Linguistics.

Rui Yan. 2016. i, poet: Automatic poetry composition through recurrent neural networks with iterative polishing schema. In *IJCAI*.

Lili Yao, Nanyun Peng, Ralph Weischedel, Kevin Knight, Dongyan Zhao, and Rui Yan. 2019. Plan-and-write: Towards better automatic storytelling. In *AAAI*, pages 7378–7385.

Steve Young, Milica Gasic, Simon Keizer, Francois Mairesse, Jost Schatzmann, Blaise Thomson, and Kai Yu. 2010. The hidden information state model: A practical framework for pomdp-based spoken dialogue management. *Computer Speech and Language*, 24(2):150–174.

Rowan Zellers, Ari Holtzman, Hannah Rashkin, Yonatan Bisk, Ali Farhadi, Franziska Roesner, and Yejin Choi. 2019. Defending against neural fake news. In *NeurIPS*.

Yaoming Zhu, Sidi Lu, Lei Zheng, Jiaxian Guo, Weinan Zhang, Jun Wang, and Yong Yu. 2018. Texygen: A benchmarking platform for text generation models. *CoRR*, abs/1802.01886. https://github.com/geek-ai/Texygen.

A Supplementary Materials

A.1 Examples from Training Datasets

We show full stories in Tables 6-8 corresponding to the excerpts shown in the Dataset sub-section of Outline-Conditioned Generation in the main text.

A.2 Human Evaluation Details

In Figures 5-8, we show the questionaires we asked the human raters. In question 2 of the full story task, we asked about which story was *more* repetitive, but we flip their answers in Figure 3 to show the model that was *less* repetitive in the Figure (i.e. for ease of reading, we made higher better as with the other metrics).

Q1: How repetitive is the information in this paragraph of information from the previous paragraph?
- **5- Very Repetitive**
- **4**
- **3- Somewhat Repetitive**
- **2**
- **1- Not Repetetive**
- gibberish

Try to explain why (in a few words or a short sentence):

Q2: How smooth/natural is the transition to this paragraph from the previous paragraph?
- **5- Smooth Transition**
- **4**
- **3- Decent Transition**
- **2**
- **1- Abrupt/Awkward Transition**
- gibberish

Try to explain why (in a few words or a short sentence):

Q3: How much does this paragraph follow a single plotline?
- **5- Consistently**
- **4**
- **3- Somewhat**
- **2**
- **1- Not at all**
- gibberish

Try to explain why (in a few words or a short sentence):

Figure 5: Questionaire for the narrative flow questions about paragraph excerpts. We pay humans $1.00 per HIT.

Q1: Which do you think is the correct order of these two paragraphs?
- **Paragraph 1, then Paragraph 2**
- **Paragraph 2, then Paragraph 1**
- **Either is fine**
- these paragraphs are too gibberish to be understood

Q2: Try to explain why (in a few words or a short sentence):

Figure 6: Questionaire for the ordering questions about paragraph excerpts. We pay humans $1.00 per HIT.

Q1: Which do you think is better at utilizing the keywords?
- **Paragraph 1**
- **Paragraph 2**

Q2: Try to explain why (in a few words or a short sentence):

Figure 7: Questionaire for the head-to-head outline usage questions about paragraph excerpts. We pay humans $1.00 per HIT.

Q1: Which do you think is better at utilizing the keywords?
- **Story 1**
- **Story 2**

Q2: Which do you think is more repetitive?
- **Story 1**
- **Story 2**

Q3: Which do you think has better transitions?
- **Story 1**
- **Story 2**

Q4: Which do you think is better at following a single storyline?
- **Story 1**
- **Story 2**

Q5: Which do you think has a better introduction?
- **Story 1**
- **Story 2**

Q6: Which do you think has a better conclusion?
- **Story 1**
- **Story 2**

Q7: Which do you think has a clear order of events?
- **Story 1**
- **Story 2**

Figure 8: Questionaire for the head-to-head questions about full stories. We pay humans $2.00 per HIT. Note: we reversed the answers to question 2 so that we could show which models were *less* repetitive in Figure 3.